

Controls before autonomy

An AI agent with real access needs the safeguards a new employee gets on day one. Governance decides what it is allowed to do. **Controls are how it stays inside that line once it is live:** limited access, a record of what it did, a way to undo it, someone watching, and a person checking its early work.



THE FIVE CONTROLS ————— each one is a thing you'd give a new hire

- I

Access LEAST PRIVILEGE

The narrowest access that lets it do the job, and nothing beyond. If it only needs to read, it should never be able to write. **Wider access only means a mistake can reach further.**
- 2

Audit trail A RECORD OF WHAT IT DID

Every meaningful action logged in plain language, so on Monday you can answer what it did and when. **A log only helps if someone can read it.**
- 3

Rollback THE UNDO BUTTON

Know how you'd undo an action before you grant the access. If you can't cleanly reverse it, it probably shouldn't be automated yet.
- 4

Monitoring SOMEONE WOULD NOTICE

Alerts you set on purpose: a metric moving too far, too many actions too fast, a decision it's unsure about. **No alerts often means nobody is watching.**
- 5

Human-in-the-loop THE PROBATION

Start with close review, then give it more independence as it earns trust. Never the other way round.

HOW MUCH INDEPENDENCE: THE AUTONOMY DIAL ————— one level at a time



The level you grant must match the controls you've actually built. **The common mistake is level-3 autonomy on level-0 controls:** no real audit, no rollback, nothing watching. Start low, and move up one level at a time as the controls underneath are built.

Monday morning YOUR FIRST MOVE

Take one AI agent you already run – or one you're about to switch on – and answer the five onboarding questions **out loud**: what exactly can it touch, where's the log of what it did, how would you undo its last action, what would set off an alert, and who is actually reviewing it right now?

Whichever one you can't answer cleanly is the control to fix first. **Fix it before more autonomy, and before the next use case.**

OPERATOR INSIGHT · LAYER 4

Controls aren't bureaucracy - they're operational safety systems. So a fast mistake doesn't quietly become a slow disaster.

"Your AI maturity will never exceed your operational maturity."

Free

Score your own controls – the AI-Ready Operations Assessment runs all five layers as a readiness check, on a scale from "ad hoc" to "AI-ready". [vihrenlabs.gumroad.com/l/rmkje](https://vihrenlabs.com/l/rmkje)